
MODULE-WISE COURSE OVERVIEW (30–50 HOUR AI/ML COURSE)

1. INTRODUCTION TO ARTIFICIAL INTELLIGENCE

- WHAT IS AI? DEFINITION AND HISTORY
 - TYPES OF AI: NARROW, GENERAL, SUPERINTELLIGENT
 - REAL-WORLD APPLICATIONS
 - OVERVIEW OF ML, DL, NLP, CV, GENAI
 - MODERN AI ECOSYSTEM: LLMs, MULTIMODAL AI
-

2. MACHINE LEARNING (ML)

- WHAT IS ML?
 - SUPERVISED, UNSUPERVISED, REINFORCEMENT LEARNING
 - DATA PREPROCESSING, FEATURE ENGINEERING, EDA
 - REGRESSION & CLASSIFICATION MODELS
 - ENSEMBLE METHODS: BAGGING, BOOSTING
 - CLUSTERING: K-MEANS, DBSCAN, HIERARCHICAL
 - DIMENSIONALITY REDUCTION: PCA, LDA, t-SNE
 - EVALUATION & TUNING: GRID/RANDOM SEARCH, REGULARIZATION, OVERFITTING, BIAS–VARIANCE
-

3. DEEP LEARNING (DL)

- INTRODUCTION TO NEURAL NETWORK
- PERCEPTRON, ACTIVATIONS, BACKPROPAGATION
- FULLY CONNECTED NETWORKS
- CNNs
- RNN, LSTM, GRU

- GNN
 - TRANSFORMERS & ATTENTION
 - GANs vs DIFFUSION MODELS
 - LOSS FUNCTIONS, OPTIMIZERS, REGULARIZATION
-

4. NATURAL LANGUAGE PROCESSING (NLP)

- CLASSICAL NLP: TOKENIZATION, STOPWORDS, STEMMING, TF-IDF
 - TEXT EMBEDDINGS (WORD2VEC, GLOVE), TOKENIZATION,
 - RNN/LSTM/GRU
 - ATTENTION
 - CONTEXT
 - TRANSFORMERS: BERT, GPT, T5
 - TASKS: NER, QA, SUMMARIZATION, TRANSLATION
 - PROMPT ENGINEERING
 - TOOLS: NLTK, SPACY, HUGGINGFACE
 - NLP PROJECTS
-

5. COMPUTER VISION (CV)

- COMPUTER VISION BASICS AND CHALLENGES
- IMAGE FORMATS, PIXELS
- CONVOLUTIONS, FILTERS, SIFT, HOG
- OBJECT DETECTION: R-CNN → YOLO → DETR
- SEGMENTATION: U-NET, MASK R-CNN
- FACE DETECTION/RECOGNITION
- VISION TRANSFORMER (ViT)
- CV PROJECTS

6. TRANSFER LEARNING & FINE-TUNING

- WHY PRETRAINED MODELS?
 - REUSING MODELS IN NLP, CV, MULTIMODAL TASKS
 - FINE-TUNING METHODS: FULL, ADAPTERS, LoRA, QLoRA,...
 - PROMPT-TUNING VS FINE-TUNING
 - USE CASES
-

7. RETRIEVAL-AUGMENTED GENERATION (RAG)

- WHY RAG?
 - RAG PIPELINE (EMBED → INDEX → RETRIEVE → GENERATE)
 - VECTOR DATABASES (FAISS, CHROMA, MILVUS)
 - CHUNKING STRATEGIES
 - RETRIEVAL EVALUATION
 - RAG VS FINE-TUNING
-

8. SPEECH AI

- ASR (WHISPER)
- TTS (VITS, BARK)
- SPEAKER IDENTIFICATION
- EMOTIONrecognition
- VOICE ASSISTANTS
- SPEECH + LLM MULTIMODAL AGENTS

9. GRAPH NEURAL NETWORKS (GNNs)

OVERVIEW

- WHAT ARE GRAPH NEURAL NETWORKS AND WHY DO WE NEED THEM?
- REPRESENTING DATA AS GRAPHS (NODES, EDGES, ADJACENCY MATRICES)
- GRAPH LEARNING VS TRADITIONAL ML/DL
- APPLICATIONS: SOCIAL NETWORKS, RECOMMENDATION SYSTEMS, FRAUD DETECTION, MOLECULAR PREDICTIONS, KNOWLEDGE GRAPHS

CORE CONCEPTS

- GRAPH CONVOLUTION
- MESSAGE PASSING
- NEIGHBORHOOD AGGREGATION
- GRAPH EMBEDDINGS

POPULAR GNN ARCHITECTURES

- GCN (GRAPH CONVOLUTIONAL NETWORK)
- GAT (GRAPH ATTENTION NETWORK)
- GraphSAGE
- RELATIONAL GCN (R-GCN)
- HETEROGENEOUS GRAPH MODELS

ADVANCED GNN TOPICS

- LINK PREDICTION
- NODE CLASSIFICATION
- GRAPH CLASSIFICATION
- KNOWLEDGE GRAPH EMBEDDINGS
- COMBINING LLMs + GNNs (EMERGING FIELD)

TOOLS

- PYTORCH GEOMETRIC (PyG)
- DGL (DEEP GRAPH LIBRARY)

GNN PROJECTS

- SOCIAL NETWORK NODE CLASSIFICATION

- MOLECULAR PROPERTY PREDICTION
 - FRAUD DETECTION USING TRANSACTION GRAPHS
-

10. ADVANCED TOPICS

- DETR
 - MULTIMODALITY
 - DECODING STRATEGIES
 - MoE
 - LARGE LANGUAGE MODELS: ARCHITECTURE, TRAINING, APPLICATIONS (CHATGPT, LLAMA, MISTRAL)
-

11. VISION–LANGUAGE MODELS (VLMs)

- WHAT ARE VLMs?
 - CLIP
 - GROUNDINGDINO (LANGUAGE-GUIDED OBJECT DETECTION)
 - SAM (FOUNDATION SEGMENTATION MODEL)
 - GROUNDED-SAM WORKFLOW
 - APPLICATIONS: MULTIMODAL SEARCH, ROBOTICS, VQA
-

12. AI EXPLAINABILITY

- EXPLAINABILITY: GRAD-CAM, ATTENTION VIZ

13. END-TO-END PROJECTS

- ML PIPELINE
- CNN CLASSIFIER
- NLP (NER, SENTIMENT)
- TRANSFORMER-BASED CHATBOT
- RAG-BASED SYSTEM
- GROUNDINGDINO + SAM DETECTOR
- WHISPER ASR TOOL
- LLM LoRA/QLoRA FINE-TUNING